

Data Wrangling with R's tidyverse

dplyr, tidyr and friends

w. cools

key message on data manipulation	3
R's tidyverse packages	4
set up tidyverse packages	4
tibbles and pipes	5
example, getting ahead of ourselves	6
dplyr to manipulate data	7
grouping	8
filter	9
select	11
mutate	15
summarize	18
scoping a verb	19
join	21
dplyr exercises	22
example, getting ahead of ourselves again	23
friends of dplyr	25
tidyr to tidy data	26
pivot	26
separate / unite	27
import data with readr, readxl or haven	29
readr	29
readxl	29
haven	30
final summary on Data Manipulation	31

Compiled Jul 12, 2020

Current draft aims to introduce researchers to data manipulation in R with the `dplyr`, `tidyr`, and `stringr` packages of the **tidyverse** ecosystem.

Our target audience is primarily the research community at VUB / UZ Brussel, those who have some basic experience in R and want to know more.

We invite you to help improve this document by sending us feedback
wilfried.cools@vub.be or anonymously at icds.be/consulting (right side, bottom)

key message on data manipulation

data manipulation is inherent to data analysis, not just a precursor

- no -fit's all data representation-, dependent on analysis or visualization (note: raw data should be unaltered)
- flexible use of data manipulation elicits better data exploration and modeling

data manipulation is best done with coding (as opposed to manual changes), provides the best guarantee to

- efficiently and correctly process data and statistics
- maintain structure and transparency, to support reproducibility

data manipulation is easier and more intuitive when maintaining tidy data

- tidy data: meaning appropriately mapped into structure
 - each row an observation as research unit,
 - each column a variable as property,
 - each cell a particular value, linking row to column
 - note: data can be split into multiple tables (relational data).
- aim for tidy data registration (avoid tedious manipulations)




Figure 1: workflow: tidyverse lingo

R's tidyverse packages

focus in current draft is on R

- free, hugely flexible, large community online to help out
- offers general functionality (base R module)
- offers many packages with dedicated functionality

in particular on the **tidyverse** package (Hadley Wickham et al.), an ecosystem that includes

- **dplyr** for manipulating data frames [main focus]
- **tidyr** for tidying data [check Data Representation]
- **stringr** for dealing with texts
- **readr** for reading in data [highlighted]
- **tibble** for data representation [highlighted]
- **forcats** for dealing with factors
- **ggplot** for visualizing data [separate draft]
- **purrr** for functional programming (advanced)
- ...

find convenient cheat sheets at <https://rstudio.com/resources/cheatsheets/>

tidyverse is a very good extension of base R, because it

- is much more consistent (functions and packages) → eco-system
- avoids poor historical choices, sets good defaults
- explicitly links to tidy data

set up tidyverse packages

Install (at least once) and load (once per R session) the **tidyverse** package.

```
install.packages('tidyverse')
```

The individual packages that are loaded are listed, as are their conflicts.

```
library(tidyverse)
```

Conflicts can arise when loading packages with the same function names. Resolve conflicts for example with explicitly referencing the package with `::`, eg., `stat::filter()`.

Conflicts can be checked for tidyverse.

```
tidyverse_conflicts( )
```

```
| -- Conflicts -----  
| x dplyr::filter() masks stats::filter()  
| x dplyr::lag()    masks stats::lag()
```

The **tidyverse** ecosystem includes

broom, cli, crayon, dbplyr, dplyr, forcats, ggplot2, haven, hms, httr, jsonlite, lubridate, magrittr, modelr, pillar, purrr, readr, readxl, reprex, rlang, rstudioapi, rvest, stringr, tibble, tidyr, xml2, tidyverse.

```
tidyverse_packages( )
```

Most of these packages should be loaded explicitly (not included in `library(tidyverse)`).

tibbles and pipes

data type: **tibble** package

- R data type for analysis is a **data.frame**, a list of equally sized vectors.
 - numeric vector (either double, integer, or complex)
 - factor (ordered, not ordered)
 - boolean vector
 - character
- a **tibble** is a **data.frame**, enhanced for convenience and consistency.
 - creating a **tibble** with the `tribble()` function (`class()` shows both **data.frame** and **tbl_df**)

```
(mytibble <- tribble(
  ~colA, ~colB,
  "a", 1,
  "b", 2,
  "c", 3
))
```

```
| # A tibble: 3 x 2
|   colA   colB
|   <chr> <dbl>
| 1 a         1
| 2 b         2
| 3 c         3
```

```
class(mytibble)
```

```
| [1] "tbl_df"      "tbl"        "data.frame"
```

creating a **data.frame**, it looks (not shown) slightly different.

```
mydf <- data.frame(colA=c('a', 'b', 'c'), colB=1:3)
class(mydf)
```

```
| [1] "data.frame"
```

process: **magrittr** package for pipes

- R processes data with functions, eg., `mean(mytibble$colB)`
 - reads inside-out, standard in base R and optional in tidyverse
- R data can be processed using pipes too, eg., `mytibble %>% summarize(mean(colB))`
 - reads left to right, standard in tidyverse
 - especially of interest with multiple steps, serves readability

example: calculate the square rooted sum of squared differences between two variables.

```
x1 <- rnorm(10); x2 <- rnorm(10)
sqrt(sum((x1-x2)^2))
```

```
| [1] 2.676459
```

```
(x1-x2)^2 %>% sum( ) %>% sqrt( )
```

```
| [1] 2.676459
```

example, getting ahead of ourselves

exemplary data (part of base R): `mtcars`.

load data available in R packages with the `data()` function, observe it's structure with the `str()` function and observe the first 6 observations with the `head()` function.

```
data(mtcars)
str(mtcars)
```

```
| 'data.frame': 32 obs. of 11 variables:
| $ mpg : num 21 21 22.8 21.4 18.7 18.1 14.3 24.4 22.8 19.2 ...
| $ cyl : num 6 6 4 6 8 6 8 4 4 6 ...
| $ disp: num 160 160 108 258 360 ...
| $ hp : num 110 110 93 110 175 105 245 62 95 123 ...
| $ drat: num 3.9 3.9 3.85 3.08 3.15 2.76 3.21 3.69 3.92 3.92 ...
| $ wt : num 2.62 2.88 2.32 3.21 3.44 ...
| $ qsec: num 16.5 17 18.6 19.4 17 ...
| $ vs : num 0 0 1 1 0 1 0 1 1 1 ...
| $ am : num 1 1 1 0 0 0 0 0 0 0 ...
| $ gear: num 4 4 4 3 3 3 3 4 4 4 ...
| $ carb: num 4 4 1 1 2 1 4 2 2 4 ...
```

```
head(mtcars)
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	160	110	3.90	2.620	16.46	0	1	4	4
Mazda RX4 Wag	21.0	6	160	110	3.90	2.875	17.02	0	1	4	4
Datsun 710	22.8	4	108	93	3.85	2.320	18.61	1	1	4	1
Hornet 4 Drive	21.4	6	258	110	3.08	3.215	19.44	1	0	3	1
Hornet Sportabout	18.7	8	360	175	3.15	3.440	17.02	0	0	3	2
Valiant	18.1	6	225	105	2.76	3.460	20.22	1	0	3	1

note: calling `data()` without arguments shows all the available data currently in reach.

a tidyverse look at the data, an alternative to `str()` is `glimpse()`.

```
glimpse(mtcars)
```

```
| Observations: 32
| Variables: 11
| $ mpg <dbl> 21.0, 21.0, 22.8, 21.4, 18.7, 18.1, 14.3, 24.4, 22.8, 19.2, 17.8, 16.4, 17.3, 15.2, 10.4
| $ cyl <dbl> 6, 6, 4, 6, 8, 6, 8, 4, 4, 6, 6, 8, 8, 8, 8, 8, 4, 4, 4, 4, 8, 8, 8, 8, 4, 4, 4, 8, 6
| $ disp <dbl> 160.0, 160.0, 108.0, 258.0, 360.0, 225.0, 360.0, 146.7, 140.8, 167.6, 167.6, 275.8, 275.8,
```

```

| $ hp    <dbl> 110, 110, 93, 110, 175, 105, 245, 62, 95, 123, 123, 180, 180, 180, 205, 215, 230, 66, 52
| $ drat  <dbl> 3.90, 3.90, 3.85, 3.08, 3.15, 2.76, 3.21, 3.69, 3.92, 3.92, 3.92, 3.07, 3.07, 3.07, 2.93
| $ wt    <dbl> 2.620, 2.875, 2.320, 3.215, 3.440, 3.460, 3.570, 3.190, 3.150, 3.440, 3.440, 4.070, 3.73
| $ qsec  <dbl> 16.46, 17.02, 18.61, 19.44, 17.02, 20.22, 15.84, 20.00, 22.90, 18.30, 18.90, 17.40, 17.6
| $ vs    <dbl> 0, 0, 1, 1, 0, 1, 0, 1, 1, 1, 1, 0, 0, 0, 0, 0, 1, 1, 1, 1, 0, 0, 0, 0, 1, 0, 1, 0, 0
| $ am    <dbl> 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1
| $ gear  <dbl> 4, 4, 4, 3, 3, 3, 3, 4, 4, 4, 4, 3, 3, 3, 3, 3, 3, 4, 4, 4, 3, 3, 3, 3, 3, 4, 5, 5, 5, 5
| $ carb  <dbl> 4, 4, 1, 1, 2, 1, 4, 2, 2, 4, 4, 3, 3, 3, 4, 4, 4, 1, 2, 1, 1, 2, 2, 4, 2, 1, 2, 2, 4, 6

```

an example !

take the `mtcars` data,
 select specific variables (`mpg,cyl,hp,am`) and rename one (`hp` turns to `hpow`),
 select specific rows (`mpg` bigger than 15),
 create a new variable based on existing variables (`mpgr` is the ratio `mpg` on `hpow`),
 summarize that new variable per group formed by combining two variables (minimum of `mpgr` per `cyl/am` group), and reshape the result into a table with one row per `cyl`-value (4,6,8) and a column for each `am` value (0,1), with column variable names renamed to `am0` and `am1`.

```

mtcars %>%
  select(mpg, cyl, hpow=hp, am) %>%
  filter(mpg > 15) %>%
  mutate(mpgr = mpg/hpow) %>%
  group_by(cyl, am) %>%
  summarize(min=min(mpgr)) %>%
  pivot_wider(names_from=am, values_from=min) %>%
  select(cyl, am0=`0`, am1=`1`)

```

```

| # A tibble: 3 x 3
| # Groups:   cyl [3]
|   cyl    am0    am1
|   <dbl> <dbl> <dbl>
| 1     4 0.222 0.196
| 2     6 0.145 0.113
| 3     8 0.0844 0.0598

```

dplyr to manipulate data

dplyr (package within tidyverse) is focused on

- manipulating dataframes (tibbles): subsetting, altering, summarizing, ordering, combining, reshaping
- to explore and transform (for visualization / modeling)
- applying functions to dataframes (tibbles)

main -verbs- (see example above)

- `filter()` : conditional selection of cases
- `select()` : conditional selection of variables, allows reordering and renaming
- `mutate()` : creation of new variables based on existing variables

- summarise() : reduce sets of values to single values

verb to structure data (see example above)

- group_by() : internal grouping, undo with ungroup()
- works preceding main verbs

verbs enhanced with control on scope (advanced)

- across() : new way of scoping (instead of *_it, *_at, *_all)
- works for selection in mutate() and summarize()

additional dplyr verbs:

- arrange() : ordering of cases
- sample_n() and sample_frac() : random sampling
- slice(), transmute(), rename(), relocate(), ...

verbs to extend data

- bind_rows() and bind_cols() : append data of same structure
- left_, right_, inner_, full_, semi_ and anti_ join() : join data using indicator variable(s)

final comment:

only the core of dplyr is discussed, much more is possible and you will find on the Net.

grouping

grouping prepares data for group specific operations.

a glimpse of the data shows variables and..

- number of observations and variables
- number of groups and grouping variables

```
tst <- mtcars %>% group_by(am,vs)
glimpse(tst)
```

```
| Observations: 32
| Variables: 11
| Groups: am, vs [4]
| $ mpg   <dbl> 21.0, 21.0, 22.8, 21.4, 18.7, 18.1, 14.3, 24.4, 22.8, 19.2, 17.8, 16.4, 17.3, 15.2, 10.4
| $ cyl   <dbl> 6, 6, 4, 6, 8, 6, 8, 4, 4, 6, 6, 8, 8, 8, 8, 8, 4, 4, 4, 4, 8, 8, 8, 8, 4, 4, 4, 8, 6
| $ disp  <dbl> 160.0, 160.0, 108.0, 258.0, 360.0, 225.0, 360.0, 146.7, 140.8, 167.6, 167.6, 275.8, 275.8
| $ hp    <dbl> 110, 110, 93, 110, 175, 105, 245, 62, 95, 123, 123, 180, 180, 180, 205, 215, 230, 66, 52
| $ drat  <dbl> 3.90, 3.90, 3.85, 3.08, 3.15, 2.76, 3.21, 3.69, 3.92, 3.92, 3.92, 3.07, 3.07, 3.07, 2.93
| $ wt    <dbl> 2.620, 2.875, 2.320, 3.215, 3.440, 3.460, 3.570, 3.190, 3.150, 3.440, 3.440, 4.070, 3.730
| $ qsec  <dbl> 16.46, 17.02, 18.61, 19.44, 17.02, 20.22, 15.84, 20.00, 22.90, 18.30, 18.90, 17.40, 17.60
| $ vs    <dbl> 0, 0, 1, 1, 0, 1, 0, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 0, 0, 0, 0, 1, 0, 1, 0, 0
```



```
| $ am    <dbl> 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1
| $ gear <dbl> 4, 4, 4, 3, 3, 3, 3, 3, 4, 4, 4, 4, 3, 3, 3, 3, 3, 3, 4, 4, 4, 3, 3, 3, 3, 3, 4, 5, 5, 5, 5
| $ carb <dbl> 4, 4, 1, 1, 2, 1, 4, 2, 2, 4, 4, 3, 3, 3, 4, 4, 4, 1, 2, 1, 1, 2, 2, 4, 2, 1, 2, 2, 4, 6, 6
```

```
tst %>% summarize(n( ))
```

```
| # A tibble: 4 x 3
| # Groups:   am [2]
|   am    vs `n()`
|   <dbl> <dbl> <int>
| 1     0     0    12
| 2     0     1     7
| 3     1     0     6
| 4     1     1     7
```

actions on grouped data are grouped too, eg., count the number of observations (`n()`).

remove grouping with `ungroup()`, good practice to avoid side effects.

```
tst <- tst %>% ungroup( )
tst %>% summarize(n( ))
```

```
| # A tibble: 1 x 1
|   `n()`
|   <int>
| 1    32
```

help file shows ways for consecutive grouping with `.add` and `.drop` arguments.

transformed variable can be used for grouping, for example cutting the `mpg` in 3 groups.

```
tst <- mtcars %>% group_by(mpg3 = cut(mpg, 3))
tst %>% summarize(n( ))
```

```
| # A tibble: 3 x 2
|   mpg3    `n()`
|   <fct>    <int>
| 1 (10.4,18.2]    14
| 2 (18.2,26.1]    13
| 3 (26.1,33.9]     5
```

filter

Return rows using matching conditions.

example: the `mpg` could be set with a minimum of 30.

```
mtcars %>% filter(mpg > 30)
```

```
|           mpg cyl disp  hp drat   wt  qsec vs am gear carb
| Fiat 128    32.4   4  78.7  66 4.08 2.200 19.47  1  1    4    1
```

```
| Honda Civic      30.4   4 75.7  52 4.93 1.615 18.52  1  1   4   2
| Toyota Corolla  33.9   4 71.1  65 4.22 1.835 19.90  1  1   4   1
| Lotus Europa    30.4   4 95.1 113 3.77 1.513 16.90  1  1   5   2
```

More than one condition can be considered jointly.

example: extract rows with mpg above 20 AND qsec below or equal to 18.

```
mtcars %>% filter(mpg > 20, qsec <= 18)
```

```
|           mpg cyl  disp  hp drat   wt  qsec vs am gear carb
| Mazda RX4      21.0   6 160.0 110 3.90 2.620 16.46  0  1   4   4
| Mazda RX4 Wag  21.0   6 160.0 110 3.90 2.875 17.02  0  1   4   4
| Porsche 914-2  26.0   4 120.3  91 4.43 2.140 16.70  0  1   5   2
| Lotus Europa   30.4   4  95.1 113 3.77 1.513 16.90  1  1   5   2
```

example: extract rows with mpg above 30 OR qsec below 20 AND am equal to 0.

```
mtcars %>% filter(mpg > 30 | qsec > 20, am==0)
```

```
|           mpg cyl  disp  hp drat   wt  qsec vs am gear carb
| Valiant        18.1   6 225.0 105 2.76 3.460 20.22  1  0   3   1
| Merc 230        22.8   4 140.8  95 3.92 3.150 22.90  1  0   4   2
| Toyota Corona  21.5   4 120.1  97 3.70 2.465 20.01  1  0   3   1
```

Technically it is possible to use it on grouped data.

example: calculate group average per level of vs for disp and extract only rows with values above the 90th percentile.

```
mtcars %>% group_by(vs) %>% filter(displ >= quantile(displ,.9))
```

```
| # A tibble: 4 x 11
| # Groups:   vs [2]
|   mpg   cyl  displ    hp  drat    wt  qsec    vs    am  gear  carb
|   <dbl> <dbl> <dbl>   <dbl> <dbl>   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
| 1  21.4     6   258    110  3.08  3.22  19.4     1     0     3     1
| 2  18.1     6   225    105  2.76  3.46  20.2     1     0     3     1
| 3  10.4     8   472    205  2.93  5.25  18.0     0     0     3     4
| 4  10.4     8   460    215   3     5.42  17.8     0     0     3     4
```

example: get the first 2 rows per group of vs.

```
mtcars %>% group_by(vs) %>% top_n(2)
```

```
| # A tibble: 4 x 11
| # Groups:   vs [2]
|   mpg   cyl  displ    hp  drat    wt  qsec    vs    am  gear  carb
|   <dbl> <dbl> <dbl>   <dbl> <dbl>   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
| 1  19.2     6  168.    123  3.92  3.44  18.3     1     0     4     4
| 2  17.8     6  168.    123  3.92  3.44  18.9     1     0     4     4
| 3  19.7     6  145.    175  3.62  2.77  15.5     0     1     5     6
| 4  15      8  301.    335  3.54  3.57  14.6     0     1     5     8
```

```
mtcars %>% group_by(gear) %>% distinct(cyl)
```

```
| # A tibble: 8 x 2
| # Groups:   gear [3]
|   cyl gear
|   <dbl> <dbl>
| 1     6   4
```

2	4	4
3	6	3
4	8	3
5	4	3
6	4	5
7	8	5
8	6	5

exercises on filter

- the starwars dataset is part of tidyverse, load it in !
- have a glimpse at the data, what do you see ?
- filter the data to retain only characters with light skin and brown eye color.
- arrange the data according to the character's height, largest on top !
- who is smallest ?
- slice the data and keep only the 5th to 10th observation !
- slice the top 2 observations (check help on `slice_head()`), for each gender (group your data) !
- what other functions are discussed at `?slice_head` ?
- use `slice_sample()` to randomly select 5 observations !
- use `slice_max()` to select 3 observations with highest values on `height` !
- repeat the above, but ignored characters with missing data for `mass` and get the top 3 for each species !

select

Extract columns (variables) by name, rename and/or reorder them.

example: the `mpg` could be selected.

```
mtcars %>% select(mpg)
```

	mpg
Mazda RX4	21.0
Mazda RX4 Wag	21.0
Datsun 710	22.8
Hornet 4 Drive	21.4
Hornet Sportabout	18.7
Valiant	18.1
Duster 360	14.3
Merc 240D	24.4
Merc 230	22.8
Merc 280	19.2
Merc 280C	17.8
Merc 450SE	16.4
Merc 450SL	17.3
Merc 450SLC	15.2
Cadillac Fleetwood	10.4
Lincoln Continental	10.4
Chrysler Imperial	14.7
Fiat 128	32.4

```
| Honda Civic      30.4
| Toyota Corolla   33.9
| Toyota Corona    21.5
| Dodge Challenger 15.5
| AMC Javelin      15.2
| Camaro Z28       13.3
| Pontiac Firebird 19.2
| Fiat X1-9        27.3
| Porsche 914-2    26.0
| Lotus Europa     30.4
| Ford Pantera L   15.8
| Ferrari Dino     19.7
| Maserati Bora    15.0
| Volvo 142E       21.4
```

```
# mtcars$mpg
```

Notice that even with one column, the result remains a dataframe (not a vector)

The `dplyr` way to get the base R result is by using `pull()`.

example: pull out the mpg.

```
mtcars %>% pull(mpg)
```

```
| [1] 21.0 21.0 22.8 21.4 18.7 18.1 14.3 24.4 22.8 19.2 17.8 16.4 17.3 15.2 10.4 10.4 14.7 32.4 30.4 3
| [31] 15.0 21.4
```

More than one column can be considered jointly, their order is specified as such.

example: extract columns `qsec` and `mpg` (top 6 observations).

```
mtcars %>% select(qsec,mpg) %>% head( )
```

```
|           qsec  mpg
| Mazda RX4    16.46 21.0
| Mazda RX4 Wag 17.02 21.0
| Datsun 710    18.61 22.8
| Hornet 4 Drive 19.44 21.4
| Hornet Sportabout 17.02 18.7
| Valiant      20.22 18.1
```

Columns can be extracted by their position.

example: extract third and first column.

```
mtcars %>% select(3,1)
```

```
|           disp  mpg
| Mazda RX4    160.0 21.0
| Mazda RX4 Wag 160.0 21.0
| Datsun 710    108.0 22.8
| Hornet 4 Drive 258.0 21.4
| Hornet Sportabout 360.0 18.7
| Valiant      225.0 18.1
| Duster 360    360.0 14.3
| Merc 240D     146.7 24.4
| Merc 230      140.8 22.8
| Merc 280      167.6 19.2
| Merc 280C     167.6 17.8
| Merc 450SE    275.8 16.4
```

Merc 450SL	275.8	17.3
Merc 450SLC	275.8	15.2
Cadillac Fleetwood	472.0	10.4
Lincoln Continental	460.0	10.4
Chrysler Imperial	440.0	14.7
Fiat 128	78.7	32.4
Honda Civic	75.7	30.4
Toyota Corolla	71.1	33.9
Toyota Corona	120.1	21.5
Dodge Challenger	318.0	15.5
AMC Javelin	304.0	15.2
Camaro Z28	350.0	13.3
Pontiac Firebird	400.0	19.2
Fiat X1-9	79.0	27.3
Porsche 914-2	120.3	26.0
Lotus Europa	95.1	30.4
Ford Pantera L	351.0	15.8
Ferrari Dino	145.0	19.7
Maserati Bora	301.0	15.0
Volvo 142E	121.0	21.4

example: remove columns at third to sixth position.

```
mtcars %>% select(-c(3:6))
```

	mpg	cyl	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	16.46	0	1	4	4
Mazda RX4 Wag	21.0	6	17.02	0	1	4	4
Datsun 710	22.8	4	18.61	1	1	4	1
Hornet 4 Drive	21.4	6	19.44	1	0	3	1
Hornet Sportabout	18.7	8	17.02	0	0	3	2
Valiant	18.1	6	20.22	1	0	3	1
Duster 360	14.3	8	15.84	0	0	3	4
Merc 240D	24.4	4	20.00	1	0	4	2
Merc 230	22.8	4	22.90	1	0	4	2
Merc 280	19.2	6	18.30	1	0	4	4
Merc 280C	17.8	6	18.90	1	0	4	4
Merc 450SE	16.4	8	17.40	0	0	3	3
Merc 450SL	17.3	8	17.60	0	0	3	3
Merc 450SLC	15.2	8	18.00	0	0	3	3
Cadillac Fleetwood	10.4	8	17.98	0	0	3	4
Lincoln Continental	10.4	8	17.82	0	0	3	4
Chrysler Imperial	14.7	8	17.42	0	0	3	4
Fiat 128	32.4	4	19.47	1	1	4	1
Honda Civic	30.4	4	18.52	1	1	4	2
Toyota Corolla	33.9	4	19.90	1	1	4	1
Toyota Corona	21.5	4	20.01	1	0	3	1
Dodge Challenger	15.5	8	16.87	0	0	3	2
AMC Javelin	15.2	8	17.30	0	0	3	2
Camaro Z28	13.3	8	15.41	0	0	3	4
Pontiac Firebird	19.2	8	17.05	0	0	3	2
Fiat X1-9	27.3	4	18.90	1	1	4	1
Porsche 914-2	26.0	4	16.70	0	1	5	2
Lotus Europa	30.4	4	16.90	1	1	5	2
Ford Pantera L	15.8	8	14.50	0	1	5	4

Ferrari Dino	19.7	6	15.50	0	1	5	6
Maserati Bora	15.0	8	14.60	0	1	5	8
Volvo 142E	21.4	4	18.60	1	1	4	2

Making use of **helper functions**, selections can be more automated.

Selections can be based on partial string matching, directly with `contains()` but also using regular expressions with `matches()`.

example: extract columns with names that include the string `ar` (show 6).

```
mtcars %>% select(contains('ar')) %>% head( )
```

		gear	carb
Mazda RX4	4	4	
Mazda RX4 Wag	4	4	
Datsun 710	4	1	
Hornet 4 Drive	3	1	
Hornet Sportabout	3	2	
Valiant	3	1	

example: extract columns with names that include the string `ar` but with at least one element before and after it (show 6).

```
mtcars %>% select(matches('.ar.')) %>% head( )
```

		carb
Mazda RX4	4	
Mazda RX4 Wag	4	
Datsun 710	1	
Hornet 4 Drive	1	
Hornet Sportabout	2	
Valiant	1	

Selection can be used to rename variables, but to avoid selection the `rename()` function may be more interesting.

example: rename the `cyl` into `cyl468` to reflect its values, same for `vs` and `am`, and select it together with `mpg` (show 6).

```
mtcars %>% select(mpg,cyl468=cyl,vs01=vs,am01=am) %>% head( )
```

		mpg	cyl468	vs01	am01
Mazda RX4	21.0	6	0	1	
Mazda RX4 Wag	21.0	6	0	1	
Datsun 710	22.8	4	1	1	
Hornet 4 Drive	21.4	6	1	0	
Hornet Sportabout	18.7	8	0	0	
Valiant	18.1	6	1	0	

example: rename the `cyl`, `vs` and `am` as before without selection (show 6).

```
mtcars %>% rename(cyl468=cyl,vs01=vs,am01=am) %>% head( )
```

		mpg	cyl468	disp	hp	drat	wt	qsec	vs01	am01	gear	carb
Mazda RX4	21.0	6	160	110	3.90	2.620	16.46	0	1	4	4	
Mazda RX4 Wag	21.0	6	160	110	3.90	2.875	17.02	0	1	4	4	
Datsun 710	22.8	4	108	93	3.85	2.320	18.61	1	1	4	1	
Hornet 4 Drive	21.4	6	258	110	3.08	3.215	19.44	1	0	3	1	
Hornet Sportabout	18.7	8	360	175	3.15	3.440	17.02	0	0	3	2	
Valiant	18.1	6	225	105	2.76	3.460	20.22	1	0	3	1	

note that a `select()` will include the grouping variables by default.
These grouping variables can be isolated too with the `group_cols()` function.

example: create a grouping by `vs` and `am`, and extract only those columns.

```
mtcars %>% group_by(vs,am) %>% select(group_cols( ))
```

```
| # A tibble: 32 x 2
| # Groups:   vs, am [4]
|       vs     am
|   <dbl> <dbl>
| 1     0     1
| 2     0     1
| 3     1     1
| 4     1     0
| 5     0     0
| 6     1     0
| 7     0     0
| 8     1     0
| 9     1     0
| 10    1     0
| # ... with 22 more rows
```

exercises on select

- the starwars dataset is probably still loaded into your workspace !
- select the columns hair, skin and eye color !
- use the `:` operator for consecutive columns !
- remove these columns instead of selecting them !
- select all columns with a name ending with color !
- use select to rename `homeworld` to `home_world` !
- do the same with the `rename()` function !
- select only the numeric variables, use `where()` and `is.numeric()` !
- select only those variables with names `height`, `mass` and/or `size`, if present, use `any_of()` !

mutate

Create new variables based on existing ones.

example: the new `mpg2` is the `mpg` value squared (show 6).

```
mtcars %>% mutate(mpg2=mpg^2) %>% head( )
```

```
|   mpg cyl disp  hp drat   wt  qsec vs am gear carb  mpg2
| 1  21.0   6  160 110 3.90 2.620 16.46  0  1   4    4 441.00
| 2  21.0   6  160 110 3.90 2.875 17.02  0  1   4    4 441.00
| 3  22.8   4  108  93 3.85 2.320 18.61  1  1   4    1 519.84
| 4  21.4   6  258 110 3.08 3.215 19.44  1  0   3    1 457.96
| 5  18.7   8  360 175 3.15 3.440 17.02  0  0   3    2 349.69
| 6  18.1   6  225 105 2.76 3.460 20.22  1  0   3    1 327.61
```

example: the `mpg` variable is overwritten with its value squared (show 6).

```
mtcars %>% mutate(mpg=mpg^2) %>% head( )
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
1	441.00	6	160	110	3.90	2.620	16.46	0	1	4	4
2	441.00	6	160	110	3.90	2.875	17.02	0	1	4	4
3	519.84	4	108	93	3.85	2.320	18.61	1	1	4	1
4	457.96	6	258	110	3.08	3.215	19.44	1	0	3	1
5	349.69	8	360	175	3.15	3.440	17.02	0	0	3	2
6	327.61	6	225	105	2.76	3.460	20.22	1	0	3	1

A new variable (column) can be created based on multiple existing variables.

example: the new NEWVAR is the mpg value multiplied by the vs value (show 6).

```
mtcars %>% mutate(NEWVAR=mpg*vs) %>% head( )
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb	NEWVAR
1	21.0	6	160	110	3.90	2.620	16.46	0	1	4	4	0.0
2	21.0	6	160	110	3.90	2.875	17.02	0	1	4	4	0.0
3	22.8	4	108	93	3.85	2.320	18.61	1	1	4	1	22.8
4	21.4	6	258	110	3.08	3.215	19.44	1	0	3	1	21.4
5	18.7	8	360	175	3.15	3.440	17.02	0	0	3	2	0.0
6	18.1	6	225	105	2.76	3.460	20.22	1	0	3	1	18.1

A new variable can be created based on a newly created variable.

example: the new NEWVAR is the mpg value multiplied by the vs value and this new variable is divided by the disp variable (show 6).

```
mtcars %>% mutate(NEWVAR=mpg*vs,NEWVAR2=NEWVAR/disp) %>% head( )
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb	NEWVAR	NEWVAR2
1	21.0	6	160	110	3.90	2.620	16.46	0	1	4	4	0.0	0.00000000
2	21.0	6	160	110	3.90	2.875	17.02	0	1	4	4	0.0	0.00000000
3	22.8	4	108	93	3.85	2.320	18.61	1	1	4	1	22.8	0.21111111
4	21.4	6	258	110	3.08	3.215	19.44	1	0	3	1	21.4	0.08294574
5	18.7	8	360	175	3.15	3.440	17.02	0	0	3	2	0.0	0.00000000
6	18.1	6	225	105	2.76	3.460	20.22	1	0	3	1	18.1	0.08044444

Newly created variables are added to the existing dataframe, but to isolate them, `transmute()` can be helpful.

```
mtcars %>% transmute(NEWVAR=mpg*vs,NEWVAR2=NEWVAR/disp) %>% head( )
```

	NEWVAR	NEWVAR2
1	0.0	0.00000000
2	0.0	0.00000000
3	22.8	0.21111111
4	21.4	0.08294574
5	0.0	0.00000000
6	18.1	0.08044444

Making use of **window functions**, mutations can be more automated.

example: add a column with the cumulative sum of mpg using `cumsum()` (show 6).

```
mtcars %>% mutate(NEWVAR=cumsum(mpg)) %>% head( )
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb	NEWVAR
1	21.0	6	160	110	3.90	2.620	16.46	0	1	4	4	21.0
2	21.0	6	160	110	3.90	2.875	17.02	0	1	4	4	42.0


```
| 3 22.8 4 108 93 3.85 2.320 18.61 1 1 4 1 64.8
| 4 21.4 6 258 110 3.08 3.215 19.44 1 0 3 1 86.2
| 5 18.7 8 360 175 3.15 3.440 17.02 0 0 3 2 104.9
| 6 18.1 6 225 105 2.76 3.460 20.22 1 0 3 1 123.0
```

example: add a column with indicator whether the mpg is between 20 and 22 (show 6).

```
mtcars %>% mutate(NEWVAR=between(mpg,20,22)) %>% head( )
```

```
|   mpg cyl disp  hp drat   wt  qsec vs am gear carb NEWVAR
| 1 21.0  6  160 110 3.90 2.620 16.46 0  1   4   4   TRUE
| 2 21.0  6  160 110 3.90 2.875 17.02 0  1   4   4   TRUE
| 3 22.8  4  108  93 3.85 2.320 18.61 1  1   4   1  FALSE
| 4 21.4  6  258 110 3.08 3.215 19.44 1  0   3   1   TRUE
| 5 18.7  8  360 175 3.15 3.440 17.02 0  0   3   2  FALSE
| 6 18.1  6  225 105 2.76 3.460 20.22 1  0   3   1  FALSE
```

example: add a row number dependent on the rank of mpg values, with the `row_number()` function. When arranged by mpg this is more clear.

```
mtcars %>% mutate(id=row_number(mpg)) %>% head( )
```

```
|   mpg cyl disp  hp drat   wt  qsec vs am gear carb id
| 1 21.0  6  160 110 3.90 2.620 16.46 0  1   4   4  19
| 2 21.0  6  160 110 3.90 2.875 17.02 0  1   4   4  20
| 3 22.8  4  108  93 3.85 2.320 18.61 1  1   4   1  24
| 4 21.4  6  258 110 3.08 3.215 19.44 1  0   3   1  21
| 5 18.7  8  360 175 3.15 3.440 17.02 0  0   3   2  15
| 6 18.1  6  225 105 2.76 3.460 20.22 1  0   3   1  14
```

```
mtcars %>% mutate(id=row_number(mpg)) %>% arrange(mpg) %>% head( )
```

```
|   mpg cyl disp  hp drat   wt  qsec vs am gear carb id
| 1 10.4  8  472 205 2.93 5.250 17.98 0  0   3   4   1
| 2 10.4  8  460 215 3.00 5.424 17.82 0  0   3   4   2
| 3 13.3  8  350 245 3.73 3.840 15.41 0  0   3   4   3
| 4 14.3  8  360 245 3.21 3.570 15.84 0  0   3   4   4
| 5 14.7  8  440 230 3.23 5.345 17.42 0  0   3   4   5
| 6 15.0  8  301 335 3.54 3.570 14.60 0  1   5   8   6
```

Grouping variables can isolate the operations.

example: create a grouping by vs and am, and mutate only those columns.

```
mtcars %>% group_by(vs,am) %>% mutate(id=row_number(mpg)) %>% head( )
```

```
| # A tibble: 6 x 12
| # Groups:   vs, am [4]
|   mpg   cyl  disp    hp  drat    wt  qsec    vs   am gear carb   id
|   <dbl> <dbl> <dbl>   <dbl> <dbl>   <dbl> <dbl>   <dbl> <dbl> <dbl> <dbl> <int>
| 1  21     6   160   110   3.9    2.62  16.5     0     1     4     4     4
| 2  21     6   160   110   3.9    2.88  17.0     0     1     4     4     5
| 3  22.8   4   108    93   3.85    2.32  18.6     1     1     4     1     2
| 4  21.4   6   258   110   3.08    3.22  19.4     1     0     3     1     4
| 5  18.7   8   360   175   3.15    3.44  17.0     0     0     3     2    11
| 6  18.1   6   225   105   2.76    3.46  20.2     1     0     3     1     2
```

Now for each combination of vs and am, there will be a 1 (first), 2 (second)... for id.

exercises on mutate

- the starwars dataset is probably still loaded into your workspace !
- create a new variable `height_m` with `height` divided by 100 !
- create the same new variable, but also define BMI as `mass / height_m` to the power 2 !
- use `transmute` to repeat the above mutation but keep only `height_m` and BMI !
- create a new variable with the z-score of height for each species ($zcore = (value - mean) / sd$) !
- create that z-score per species !
- create a gender indicator that replaces the `male` and `female` labels with `m` and `f` (use `recode()`) !
- create a gender indicator that, when sex is none changes it to the species and otherwise keeps the sex specification (use `ifelse()`)!

summarize

Reduce sets of values into their summaries, based on grouped data.

A new variable (column) is created based on an existing one by summarizing, condensing the data.

example: the mean of all `mpg` values can be obtained.

```
mtcars %>% summarize(myAverage=mean(mpg))
```

```
|   myAverage
| 1  20.09062
```

Multiple new variables can be created.

example: the mean and standard deviation of all `mpg` values can be obtained, for multiple variables.

```
mtcars %>% summarize(myAvMpg=mean(mpg),mySdMpg=sd(mpg),myAvDisp=mean(displ),mySdDisp=sd(displ))
```

```
|   myAvMpg  mySdMpg myAvDisp mySdDisp
| 1 20.09062  6.026948 230.7219 123.9387
```

Grouping variables are very natural to `summarize()`.

example: the mean of all `mpg` values can be obtained for each level of `vs`.

```
mtcars %>% group_by(vs) %>% summarize(myAverage=mean(mpg))
```

```
| # A tibble: 2 x 2
|   vs myAverage
|   <dbl>   <dbl>
| 1     0    16.6
| 2     1    24.6
```

example: the mean and standard deviation of all `mpg` values can be obtained, for multiple variables, for multiple combinations of grouping, `vs` and `am`.

```
mtcars %>% group_by(vs,am) %>% summarize(myAvMpg=mean(mpg),mySdMpg=sd(mpg),myAvDisp=mean(displ),mySdDisp=sd(displ))
```

```
| # A tibble: 4 x 6
| # Groups:   vs [2]
|   vs   am myAvMpg mySdMpg myAvDisp mySdDisp
|   <dbl> <dbl>   <dbl>   <dbl>   <dbl>   <dbl>
| 1     0     0    15.0     2.77    358.     71.8
| 2     0     1    19.8     4.01    206.     95.2
```

```
| 3      1      0      20.7      2.47      175.      49.1
| 4      1      1      28.4      4.76      89.8      18.8
```

example: the total number of observations within a group, eg., **vs**, can be obtained with `n()`, or using the special verb `count()`.

```
mtcars %>% group_by(vs) %>% count( )
```

```
| # A tibble: 2 x 2
| # Groups:   vs [2]
|   vs     n
|   <dbl> <int>
| 1     0    18
| 2     1    14
```

```
mtcars %>% group_by(vs) %>% summarize(mycount=n( ))
```

```
| # A tibble: 2 x 2
|   vs mycount
|   <dbl>   <int>
| 1     0     18
| 2     1     14
```

Making use of **summary functions**, summarizing can be more automated.

example: the number of distinct values in a vector for each combination **vs** and **am** can be obtained with `n_distinct()`, and the third number of each group with `nth()`.

```
mtcars %>% group_by(vs,am) %>% summarize(nrDist=n_distinct(mpg), `3th`=nth(mpg,3))
```

```
| # A tibble: 4 x 4
| # Groups:   vs [2]
|   vs   am nrDist `3th`
|   <dbl> <dbl>   <int> <dbl>
| 1     0     0     10  16.4
| 2     0     1      5   26
| 3     1     0      7  24.4
| 4     1     1      6  30.4
```

exercises on summarize

- the starwars dataset is probably still loaded into your workspace !
- summarize the **height** into the average height (some missing values need to be dealt with) !
- repeat to above, but group by **species** and **sex**, and include the average **mass** !

scoping a verb

The `across()` function allows for selection of variables within the `summarize()` or `mutate()` function. It will replace the earlier functions `*_at`, `*_if` and `*_all`. To select the variables, they can be explicitly named or extracted with dedicated functions.

example, the structure is asked for after transforming a selected set of variables from numeric into factors.

example part 1, after turning the **am** and **vs** variable into a factor.

```
mtcars %>% select(mpg,cyl,am,vs) %>% mutate(across(c('am','vs'),factor)) %>% str( )
```

```
| 'data.frame': 32 obs. of 4 variables:
| $ mpg: num 21 21 22.8 21.4 18.7 18.1 14.3 24.4 22.8 19.2 ...
| $ cyl: num 6 6 4 6 8 6 8 4 4 6 ...
| $ am : Factor w/ 2 levels "0","1": 2 2 2 1 1 1 1 1 1 1 ...
| $ vs : Factor w/ 2 levels "0","1": 1 1 2 2 1 2 1 2 2 2 ...
```

example part 2, after turning the consecutive variables cyl, am and vs into a factor with a : operator.

```
mtcars %>% select(mpg,cyl,am,vs) %>% mutate(across(cyl:vs,factor)) %>% str( )
```

```
| 'data.frame': 32 obs. of 4 variables:
| $ mpg: num 21 21 22.8 21.4 18.7 18.1 14.3 24.4 22.8 19.2 ...
| $ cyl: Factor w/ 3 levels "4","6","8": 2 2 1 2 3 2 3 1 1 2 ...
| $ am : Factor w/ 2 levels "0","1": 2 2 2 1 1 1 1 1 1 1 ...
| $ vs : Factor w/ 2 levels "0","1": 1 1 2 2 1 2 1 2 2 2 ...
```

example part 3, after turning the variables with names that contain ar into a factor.

```
mtcars %>% select(mpg,cyl,gear,carb) %>% mutate(across(contains("ar"),factor)) %>% str( )
```

```
| 'data.frame': 32 obs. of 4 variables:
| $ mpg : num 21 21 22.8 21.4 18.7 18.1 14.3 24.4 22.8 19.2 ...
| $ cyl : num 6 6 4 6 8 6 8 4 4 6 ...
| $ gear: Factor w/ 3 levels "3","4","5": 2 2 2 1 1 1 1 2 2 2 ...
| $ carb: Factor w/ 6 levels "1","2","3","4",...: 4 4 1 1 2 1 4 2 2 4 ...
```

example, multiple values can be obtained by specifying a list of functions, for example a median, mean and sd for the first and third variable (show only first 6).

```
descr <- list(
  md = ~median(.x, na.rm = TRUE),
  av = ~mean(.x, na.rm = TRUE),
  sd = ~sd(.x, na.rm = TRUE)
)
mtcars %>% mutate(across(c(1,3), descr)) %>% head( )
```

```
|   mpg cyl disp  hp drat   wt  qsec vs am gear carb mpg_md  mpg_av  mpg_sd disp_md disp_av disp_sd
| 1 21.0   6  160  110 3.90 2.620 16.46  0  1   4    4   19.2 20.09062 6.026948  196.3 230.7219 123.93
| 2 21.0   6  160  110 3.90 2.875 17.02  0  1   4    4   19.2 20.09062 6.026948  196.3 230.7219 123.93
| 3 22.8   4  108   93 3.85 2.320 18.61  1  1   4    1   19.2 20.09062 6.026948  196.3 230.7219 123.93
| 4 21.4   6  258  110 3.08 3.215 19.44  1  0   3    1   19.2 20.09062 6.026948  196.3 230.7219 123.93
| 5 18.7   8  360  175 3.15 3.440 17.02  0  0   3    2   19.2 20.09062 6.026948  196.3 230.7219 123.93
| 6 18.1   6  225  105 2.76 3.460 20.22  1  0   3    1   19.2 20.09062 6.026948  196.3 230.7219 123.93
```

Various functions exist, as part of the tidyselect package, eg., all_of(), where(), matches(), starts_with(), and more.

While the examples use mutate(), they are also possible with summarize().

exercises on across

- the starwars dataset is probably still loaded into your workspace !
- summarize the numeric variables into their minimum and maximum (some missing values need to be dealt with) !

join

Different datafiles can be combined into one datafile using common variables that serve as key (cfr. relational databases). Methods differ primarily in how they deal with mismatches in key variable values.

example: assume a cylinder specific datafile, `mtcyl`, with a 2 cylinder but no 8 cylinder unlike the `mtcars` (4,6,8).

```
mtcyl <- tribble(
  ~cyl, ~type,
  2, 'small',
  4, 'medium',
  6, 'large'
)
```

example: combine the `mtcars` and `mtcyl` but ignore the irrelevant `cyl` equal to 2 (not part of `mtcars`), with a `left_join()`.

```
```r
mtcars %>% left_join(mtcyl) %>% head()
```

```
 | mpg cyl disp hp drat wt qsec vs am gear carb type
 | 1 21.0 6 160 110 3.90 2.620 16.46 0 1 4 4 large
 | 2 21.0 6 160 110 3.90 2.875 17.02 0 1 4 4 large
 | 3 22.8 4 108 93 3.85 2.320 18.61 1 1 4 1 medium
 | 4 21.4 6 258 110 3.08 3.215 19.44 1 0 3 1 large
 | 5 18.7 8 360 175 3.15 3.440 17.02 0 0 3 2 <NA>
 | 6 18.1 6 225 105 2.76 3.460 20.22 1 0 3 1 large
```
```

Notice that `cyl` equal to 8 turns out missing, because it is not specified in the -right- datafile.

example: combine the `mtcars` and `mtcyl` but ignore the `cyl` equal to 8 because it lacks information on type, with a `right_join()`.

```
```r
mtcars %>% right_join(mtcyl) %>% head()
```

```
 | mpg cyl disp hp drat wt qsec vs am gear carb type
 | 1 21.0 6 160.0 110 3.90 2.620 16.46 0 1 4 4 large
 | 2 21.0 6 160.0 110 3.90 2.875 17.02 0 1 4 4 large
 | 3 22.8 4 108.0 93 3.85 2.320 18.61 1 1 4 1 medium
 | 4 21.4 6 258.0 110 3.08 3.215 19.44 1 0 3 1 large
 | 5 18.1 6 225.0 105 2.76 3.460 20.22 1 0 3 1 large
 | 6 24.4 4 146.7 62 3.69 3.190 20.00 1 0 4 2 medium
```
```

Notice that `cyl` equal to 2 is included, but turns out missing for most variables because it is not specified in the -left- datafile.

example: combine the `mtcars` and `mtcyl` for only those observations with the linking variable `cyl` in both files, with an `right_join()`.

```
mtcars %>% inner_join(mtcyl) %>% head( )
```

| | mpg | cyl | disp | hp | drat | wt | qsec | vs | am | gear | carb | type |
|---|------|-----|-------|-----|------|-------|-------|----|----|------|------|--------|
| 1 | 21.0 | 6 | 160.0 | 110 | 3.90 | 2.620 | 16.46 | 0 | 1 | 4 | 4 | large |
| 2 | 21.0 | 6 | 160.0 | 110 | 3.90 | 2.875 | 17.02 | 0 | 1 | 4 | 4 | large |
| 3 | 22.8 | 4 | 108.0 | 93 | 3.85 | 2.320 | 18.61 | 1 | 1 | 4 | 1 | medium |
| 4 | 21.4 | 6 | 258.0 | 110 | 3.08 | 3.215 | 19.44 | 1 | 0 | 3 | 1 | large |
| 5 | 18.1 | 6 | 225.0 | 105 | 2.76 | 3.460 | 20.22 | 1 | 0 | 3 | 1 | large |
| 6 | 24.4 | 4 | 146.7 | 62 | 3.69 | 3.190 | 20.00 | 1 | 0 | 4 | 2 | medium |

Notice no missing values, but some data is not included.

example: combine the `mtcars` and `mtcyl` keeping all available information, with a `full_join()` showing selected rows 1 to 3, 5, 7 and 33.

```
mtcars %>% full_join(mtcyl) %>% slice(c(1:3,5,7,33))
```

| | mpg | cyl | disp | hp | drat | wt | qsec | vs | am | gear | carb | type |
|---|------|-----|------|-----|------|-------|-------|----|----|------|------|--------|
| 1 | 21.0 | 6 | 160 | 110 | 3.90 | 2.620 | 16.46 | 0 | 1 | 4 | 4 | large |
| 2 | 21.0 | 6 | 160 | 110 | 3.90 | 2.875 | 17.02 | 0 | 1 | 4 | 4 | large |
| 3 | 22.8 | 4 | 108 | 93 | 3.85 | 2.320 | 18.61 | 1 | 1 | 4 | 1 | medium |
| 4 | 18.7 | 8 | 360 | 175 | 3.15 | 3.440 | 17.02 | 0 | 0 | 3 | 2 | <NA> |
| 5 | 14.3 | 8 | 360 | 245 | 3.21 | 3.570 | 15.84 | 0 | 0 | 3 | 4 | <NA> |
| 6 | NA | 2 | NA | NA | NA | NA | NA | NA | NA | NA | NA | small |

Other types of join exist, like `semi_join()`, `nest_join()`, `anti_join()`, which are described in the helpfile.

exercises on join

- the two mini tibbles `band_members` and `band_instruments` are probably loaded into your workspace as part of the tidyverse !
- combined the two, left/right/inner/full !
- try out the same with `semi_join()` and `anti_join()` and interpret what happens.

dplyr exercises

Compare the structure of the `mtcars` data with a glimpse at that data.

Compare a select of `mpg` with a pull of `mpg`.

Check the help file and pull out the second before last column.

Select all columns except the `am`.

Select all columns except the `am` and `vs`.

Keep only columns `mpg`, `cyl` and `disp`, but rename `mpg` to `miles_gallon`.

Insist, keep only columns `mpg`, `cyl` and `disp`, but rename `mpg` to `miles per gallon`.

Keep only the consecutive columns in between `disp` and `wt`, in addition to `mpg` as a last column, use a `::`.

Create a variable for the row names.

Change the `mpg` (miles per gallon) into `kp1` (kilometers per liter) with 1 mpg is 0.425 km/l, using `mutate()`.

Select about 10% of the observations, twice, check the help file on using `sample_frac()`.

note that the matrix way could be:

Select the 10th to 15th row, check the help file on using `slice()`.

Select the distinct combinations only, for variables `am` and `vs`.

Check the help files to determine how to keep all variables (for each first observation of that combination).

Filter the data to retain only cases with `mpg > 20` and `hp` above or equal to 110.

Filter the data to retain only the Datsun 710.

example, getting ahead of ourselves again

read in data delim pivot and separate/unite

exemplary data read into an R-workspace: `repeated.txt`, a tab-delimited text file, by copy-pasting.

```
myrepeated <- read_delim(clipboard(),delim='\t')
```

```
| # A tibble: 3 x 7
|   id    `t1 score` `t1 posit` `t2 score` `t2 posit` `t3 score` `t3 posit`
|   <chr>      <dbl> <chr>      <dbl> <chr>      <dbl> <chr>
| 1 id1          1 x          NA y          4 x
| 2 id2          2 y          3 x          NA <NA>
| 3 id3          1 x          2 y          5 x
```

To get the data tidy, the different time points should not be at different columns. (See draft on Data Representation)

take the id and scores at time points and pivot it to have all scores at their own designated row, with times named `type` take the id and positions at time points and pivot it to have all positions at their own designated row too, with times named `type` separate the time part from the type part in the type variable for both scores and positions and name the time part `time` merge both datasets with a join like before, using id as key, but first eliminate at least one of the inconsistent `type` variables select only the relevant variables, id, time, score and posit, and remove all observations with a missing value for either score or posit.

```
| # A tibble: 9 x 3
|   id    type    score
|   <chr> <chr>   <dbl>
| 1 id1  t1 score    1
| 2 id1  t2 score   NA
| 3 id1  t3 score    4
| 4 id2  t1 score    2
| 5 id2  t2 score    3
| 6 id2  t3 score   NA
| 7 id3  t1 score    1
| 8 id3  t2 score    2
| 9 id3  t3 score    5

| # A tibble: 9 x 3
|   id    type    posit
|   <chr> <chr>   <chr>
| 1 id1  t1 posit x
| 2 id1  t2 posit y
```

```

| 3 id1    t3 posit x
| 4 id2    t1 posit y
| 5 id2    t2 posit x
| 6 id2    t3 posit <NA>
| 7 id3    t1 posit x
| 8 id3    t2 posit y
| 9 id3    t3 posit x

| # A tibble: 9 x 4
|   id    time type  score
|   <chr> <chr> <chr> <dbl>
| 1 id1    t1    score    1
| 2 id1    t2    score   NA
| 3 id1    t3    score    4
| 4 id2    t1    score    2
| 5 id2    t2    score    3
| 6 id2    t3    score   NA
| 7 id3    t1    score    1
| 8 id3    t2    score    2
| 9 id3    t3    score    5

| # A tibble: 9 x 4
|   id    time type  posit
|   <chr> <chr> <chr> <chr>
| 1 id1    t1    posit x
| 2 id1    t2    posit y
| 3 id1    t3    posit x
| 4 id2    t1    posit y
| 5 id2    t2    posit x
| 6 id2    t3    posit <NA>
| 7 id3    t1    posit x
| 8 id3    t2    posit y
| 9 id3    t3    posit x

| # A tibble: 9 x 5
|   id    time score type  posit
|   <chr> <chr> <dbl> <chr> <chr>
| 1 id1    t1      1 posit x
| 2 id1    t2     NA posit y
| 3 id1    t3      4 posit x
| 4 id2    t1      2 posit y
| 5 id2    t2      3 posit x
| 6 id2    t3     NA posit <NA>
| 7 id3    t1      1 posit x
| 8 id3    t2      2 posit y
| 9 id3    t3      5 posit x

| # A tibble: 7 x 4
|   id    time score posit
|   <chr> <chr> <dbl> <chr>
| 1 id1    t1      1 x
| 2 id1    t3      4 x
| 3 id2    t1      2 y
| 4 id2    t2      3 x
| 5 id3    t1      1 x
| 6 id3    t2      2 y

```



```
| 7 id3    t3          5 x
```

```
scores <- myrepeated %>% select(id, `t1 score`, `t2 score`, `t3 score`) %>%
  pivot_longer(-id, names_to='type', values_to='score') %>%
  separate(type, c('time', 'type'))
positions <- myrepeated %>% select(id, `t1 posit`, `t2 posit`, `t3 posit`) %>%
  pivot_longer(-id, names_to='type', values_to='posit') %>%
  separate(type, c('time', 'type'))
longform <- scores %>%
  select(-type) %>%
  full_join(positions) %>%
  select(-type) %>%
  filter(!is.na(score), !is.na(posit))
```

It is possible to switch back to a wider data representation, for example to calculate correlations. Maybe fill in the missing values NA as 0 values.

```
| # A tibble: 3 x 6
|   id    t1_x t3_x t1_y t2_x t2_y
|   <chr> <dbl> <dbl> <dbl> <dbl> <dbl>
| 1 id1      1     4    NA    NA    NA
| 2 id2     NA    NA     2     3    NA
| 3 id3      1     5    NA    NA     2
```

```
longform %>% pivot_wider(values_from=score, names_from=c(time, posit))
# longform %>% pivot_wider(values_from=score, names_from=c(time, posit), values_fill=list(score=0))
```

friends of dplyr

dplyr is the main package for data transformations




Figure 2: workflow: tidyverse lingo

in preparation of data manipulation

- the data has to be brought into the R workspace
- the data has to be tidy for efficient further processing

after the transformation, the data should be ready for

- modeling
- visualization [see ggplot in Data Visualization]

tidyr to tidy data

tidyr (package within tidyverse) is focused on

- tidying dataframes (tibbles): pivoting data into longer or wider form
- to explore and transform (for visualization / modeling)
- creating pure variables

main -verbs- (see example above)

- `pivot_wider()` and `pivot_longer()`: turn multiple columns or rows into one, making datafiles longer or wider
- `separate()` and `extract()`: create multiple columns from one column using delimiters or regular expressions

pivot

turning long form data into wide form and vise verse

- each research unit is assigned to a row, in a tidy dataframe (tibble)
- the research unit in focus can change throughout an analysis (eg., test score → student)
- both univariate and multivariate data representations can be required for data analysis
- multivariate data representation most intuitive, univariate most flexible

example: for the `iris` dataset, with 4 values for each unit within each species, showing only the first 6 observations.

```
long_iris <- iris %>% pivot_longer(~Species,names_to='type',values_to='score')
long_iris %>% head( )
```

```
| # A tibble: 6 x 3
|   Species type      score
|   <fct>   <chr>    <dbl>
| 1 setosa Sepal.Length  5.1
| 2 setosa Sepal.Width   3.5
| 3 setosa Petal.Length  1.4
| 4 setosa Petal.Width   0.2
| 5 setosa Sepal.Length  4.9
| 6 setosa Sepal.Width   3
```

example: for the long form `iris` dataset, the univariate case, it can not be switched into a wider form.

```
long_iris %>% pivot_wider(values_from=score,names_from=type)
```

there should be a unique combination for rows x columns, the new dataset does not link values from an individual unit (row).

example: redo the pivoting from wide to long, after adding an indicator variable for each unit.

```
long_iris <- iris %>% mutate(id=1:n()) %>% pivot_longer(~c(Species,id),names_to='type',values_to='score')
long_iris %>% head( )
```

```
| # A tibble: 6 x 4
|   Species    id type      score
|   <fct>    <int> <chr>    <dbl>
| 1 setosa      1 Sepal.Length  5.1
| 2 setosa      1 Sepal.Width   3.5
| 3 setosa      1 Petal.Length  1.4
| 4 setosa      1 Petal.Width   0.2
| 5 setosa      2 Sepal.Length  4.9
| 6 setosa      2 Sepal.Width    3
```

Also in long form, rows are combined into clusters of scores related to the same unit.

example: the pivoting can now be done from long to wide (using id information to assign scores to the appropriate row)

```
wide_iris <- long_iris %>% pivot_wider(values_from=score,names_from=type)
wide_iris
```

```
| # A tibble: 150 x 6
|   Species    id Sepal.Length Sepal.Width Petal.Length Petal.Width
|   <fct>    <int>      <dbl>      <dbl>      <dbl>      <dbl>
| 1 setosa      1      5.1        3.5        1.4        0.2
| 2 setosa      2      4.9         3         1.4        0.2
| 3 setosa      3      4.7        3.2        1.3        0.2
| 4 setosa      4      4.6        3.1        1.5        0.2
| 5 setosa      5      5         3.6        1.4        0.2
| 6 setosa      6      5.4        3.9        1.7        0.4
| 7 setosa      7      4.6        3.4        1.4        0.3
| 8 setosa      8      5         3.4        1.5        0.2
| 9 setosa      9      4.4        2.9        1.4        0.2
| 10 setosa     10      4.9        3.1        1.5        0.1
| # ... with 140 more rows
```

Note in wider format, cluster information is implied by the row, in longer format it is made explicit with indicator variables. The original wide and long distinction is abolished because it is recognized that there are various shades in between, hence wide-r and long-er.

exercises on pivot

- the `world_bank_pop` dataset that is part of the `tidyr` package
- pivot the dataset to have univariate data for the scores over the different years
- the `us_rent_income` dataset is also part of the `tidyr` package
- pivot the dataset to have a multivariate version with variables for all estimate x moe combinations
- first verify what happens when constructing a multivariate version only for estimate

separate / unite

splitting up information within a variable, or combining over variables

- each variable should consist of one type of information, in a tidy dataframe (tibble)
- variables that combine information should often be split

- variables that provide no additional information should be removed, sometimes united

example: the long form iris data shows a type that consists of both Petal/Sepal and Length/Width, the can be separated.

```
long_iris_x <- long_iris %>% separate(type,c('PT','lw'))
long_iris_x %>% head( )
```

```
| # A tibble: 6 x 5
|   Species    id PT    lw    score
|   <fct>    <int> <chr> <chr> <dbl>
| 1 setosa      1 Sepal Length  5.1
| 2 setosa      1 Sepal Width  3.5
| 3 setosa      1 Petal Length  1.4
| 4 setosa      1 Petal Width  0.2
| 5 setosa      2 Sepal Length  4.9
| 6 setosa      2 Sepal Width   3
```

example: the separated columns can be combined, using a separator dash in this case (default is underscore).

```
unite_iris <- long_iris_x %>% unite('myType',PT:lw,sep='-')
unite_iris
```

```
| # A tibble: 600 x 4
|   Species    id myType    score
|   <fct>    <int> <chr>    <dbl>
| 1 setosa      1 Sepal-Length  5.1
| 2 setosa      1 Sepal-Width  3.5
| 3 setosa      1 Petal-Length  1.4
| 4 setosa      1 Petal-Width  0.2
| 5 setosa      2 Sepal-Length  4.9
| 6 setosa      2 Sepal-Width   3
| 7 setosa      2 Petal-Length  1.4
| 8 setosa      2 Petal-Width  0.2
| 9 setosa      3 Sepal-Length  4.7
|10 setosa      3 Sepal-Width  3.2
| # ... with 590 more rows
```

The tidyr package includes some other functions that can be of interest when getting more involved into programming and simulations studies, like `expand()`, `crossover()`, `nesting()`. Best check the helpfile.

exercises on separate / unite

- the `mtcars` should still be loaded into your workspace
- turn the rownames to a variable called `type` using the `rownames_to_column()` function, it consists of car type information, car subtype and subtype specification.
- separate the `type` into three pieces (if there are less, make third and if necessary the second a missing value)
- unite the second and third part under the name `subtype`.

import data with readr, readxl or haven

when using your own data, they have to be imported into the workspace.

- data that are saved as R objects in a workspace (*.RData) can be loaded with the `load()` function
- various packages and the R base package offer functions for various types of data (excel, spss, sas, ...)
- in tidyverse the **readr** package deals with most common data, **readxl** is dedicated to notorious excel, and **haven** addresses the link with the main statistical software SAS, spss and Stata.

readr

A common function that is of interest is the `read_delim()` function from the **readr** package. A file in the working directory can be specified jointly with the delimiter, it can be obtained from the clipboard of searched through.

```
myrepeated <- read_delim(file='repeated.txt',delim='\t') # if repeated.txt is in the working directory
myrepeated <- read_delim(clipboard(),delim='\t')
myrepeated <- read_delim(file.choose(),delim='\t')
```

The `file.choose()` and `clipboard()` can be used with other functions as well, like the ones discussed next.

Consult with `?read_delim()` on how to set different parameters and gain flexibility to read in data.

readxl

Dedicated to Excel, the `read_excel()` function facilitates reading Excel files. Alternatively, save Excel files to tab-delimited or comma separated value files to read with `read_delim()`.

Example, read in the `RealData_clean.xlsx` file making use of the defaults.

```
read_excel('RealData_clean.xlsx')
```

```
| # A tibble: 280 x 21
|   Pte  Dx      P      A `Location opera~ `Size mm (FMT,V~ `FMT (mm)` `MSD (mm)` `RCOP (mm)` `Measu
|   <chr> <chr> <dbl> <dbl>          <dbl> <chr>          <dbl>      <dbl>      <dbl> <chr>
|  1 1    TB      3      1          NA VZ 24 (vierling)      NA      24      NA Y
|  2 2    TB      1     NA          NA FMT 13              13      NA      NA Y/N
|  3 3    FLIN     1     NA          NA VZ 23              NA      23      NA Y
|  4 4    TB      1     NA          NA FMT 12              12      NA      NA Y
|  5 5    TB      2     NA          NA FMT 8               8       NA      NA N
|  6 6    TB      3      2          NA FMT 6               6       NA      NA N
|  7 7    FLIN     1      1          NA FMT 2               2       NA      NA Y
|  8 8    TB      2     NA          1 FMT 8               8       NA      NA Y
|  9 9    QRT      1      1          NA QRT 10              NA      NA      10 Y
| 10 10   TB      4     NA          NA VZ 45              NA      45      NA Y
| # ... with 270 more rows, and 8 more variables: `Delay Dx-evac` <chr>, `evac <20d` <dbl>, `evac > 20d`
| #   Pil <chr>, Age <dbl>
```

Consult with `?read_excel` on alternative parameter values to extract specific columns, specifics Excel-tabs, ...

haven

Haven is dedicated to the major statistical software packages, SPSS, SAS and STATA.

An SPSS version of the iris data is part of the tidyverse system, located in the examples folder of the haven package.

The data is simply read, using default parameters, as `read_sav`.

```
pathToIrisSpssData <- system.file("examples", "iris.sav", package = "haven")
read_sav(pathToIrisSpssData)
```

```
| # A tibble: 150 x 5
|   Sepal.Length Sepal.Width Petal.Length Petal.Width Species
|   <dbl>         <dbl>         <dbl>         <dbl> <dbl+lbl>
| 1         5.1         3.5         1.4         0.2 1 [setosa]
| 2         4.9         3         1.4         0.2 1 [setosa]
| 3         4.7         3.2         1.3         0.2 1 [setosa]
| 4         4.6         3.1         1.5         0.2 1 [setosa]
| 5          5         3.6         1.4         0.2 1 [setosa]
| 6         5.4         3.9         1.7         0.4 1 [setosa]
| 7         4.6         3.4         1.4         0.3 1 [setosa]
| 8          5         3.4         1.5         0.2 1 [setosa]
| 9         4.4         2.9         1.4         0.2 1 [setosa]
| 10        4.9         3.1         1.5         0.1 1 [setosa]
| # ... with 140 more rows
```

A SAS version of the iris data is simply read with `read_sas`. The dataset has a `sas7dat` extension.

```
path <- system.file("examples", "iris.sas7bdat", package = "haven")
read_sas(path)
```

```
| # A tibble: 150 x 5
|   Sepal_Length Sepal_Width Petal_Length Petal_Width Species
|   <dbl>         <dbl>         <dbl>         <dbl> <chr>
| 1         5.1         3.5         1.4         0.2 setosa
| 2         4.9         3         1.4         0.2 setosa
| 3         4.7         3.2         1.3         0.2 setosa
| 4         4.6         3.1         1.5         0.2 setosa
| 5          5         3.6         1.4         0.2 setosa
| 6         5.4         3.9         1.7         0.4 setosa
| 7         4.6         3.4         1.4         0.3 setosa
| 8          5         3.4         1.5         0.2 setosa
| 9         4.4         2.9         1.4         0.2 setosa
| 10        4.9         3.1         1.5         0.1 setosa
| # ... with 140 more rows
```

A Stata version of the iris data is read in with `read_dta()` or `read_stata()`. The dataset has a `dta` extension.

```
path <- system.file("examples", "iris.dta", package = "haven")
read_dta(path)
```

```
| # A tibble: 150 x 5
|   sepallength sepalwidth petallength petalwidth species
|   <dbl>         <dbl>         <dbl>         <dbl> <chr>
```

```
| 1      5.10      3.5      1.40      0.200 setosa
| 2      4.90      3      1.40      0.200 setosa
| 3      4.70      3.20     1.30      0.200 setosa
| 4      4.60      3.10     1.5      0.200 setosa
| 5      5        3.60     1.40      0.200 setosa
| 6      5.40      3.90     1.70     0.400 setosa
| 7      4.60      3.40     1.40      0.300 setosa
| 8      5        3.40     1.5      0.200 setosa
| 9      4.40      2.90     1.40      0.200 setosa
| 10     4.90      3.10     1.5      0.100 setosa
| # ... with 140 more rows
```

To write any of the files, use the `write_` prefix, for `dta`, `sas` and `sav`. example, write the `mtcars` into `sas` format.

```
write_sas(mtcars, 'mytryinSAS.sas7bdat')
```

final summary on Data Manipulation

Current draft provides a primer on data manipulation, tidy data and the importing of data, which are the main steps in preparation of most real data analyses and visualizations. It is strongly advised to play with the techniques discussed above to get some proficiency in using it, as it would add significantly to the flexibility of whatever you want to further do with your data.




Figure 3: workflow: tidyverse lingo

A different draft addresses what is tidy data, with a focus on how data should be registered. A next draft will address how to visualize data, using the `ggplot()` function.

Several tidyverse packages are not yet discussed, which does suggest they are not useful but they are more specific. The consistency within tidyverse should give you a push though, to study the other packages yourself when of interest.

Base R still is a proper alternative to the tidyverse package, so be aware that others may do things differently.